

EntropyX: PE Feature Extractor for Malware Classification

Project Goal: Build a robust feature extractor for Malware Classification using Random Forest.

Last Updated: February 2026

Table of Contents

1. Project Overview
2. The Problem & Why Entropy Matters
3. Development Timeline (Git Commit History)
4. Key Technical Decisions
5. Pythonic Implementation Details
6. Architecture & Code Organization
7. Model Training & Results
8. Future Improvements

1. Project Overview

This project extracts features from Portable Executable (PE) files to train a machine learning model that distinguishes **malware** from **benign** executables. The primary signal is **Shannon Entropy**, which measures randomness in data—encrypted or compressed malware payloads have distinctively high entropy.

Key Files

File	Purpose
extractor.py	Main feature extraction logic
builder.py	Batch processes PE files → generates dataset.csv
trainer.py	Trains Random Forest model & evaluates results

File	Purpose
generate_custom_malware.sh	Creates synthetic high-entropy malware samples
generate_benign.sh	Retrieves legitimate Windows system files
data/malicious/	Training samples (custom loaders + SGN encoded)
data/benign/	Legitimate Windows executables

2. The Problem & Why Entropy Matters

Initial Hypothesis (Flawed)

"Simply calculate the entropy of the whole file to detect malware."

Why It Failed

Issue	Explanation
False Positives	Benign installers (e.g., .msi, compressed apps) also have high entropy
Missed Detections	Malware hides encrypted payloads in <i>small sections</i> , while the rest is null padding—averaging the whole file dilutes the signal
Example	custom_loader.exe showed only 1.4 entropy at file-level, but 7.13 entropy in its .data section

The Solution

Section-Aware Analysis — Analyze each PE section individually, then summarize with statistics.

3. Development Timeline (Git Commit History)

This section maps each commit to the problem it solved and the performance/accuracy improvement gained.

Commit 1: 1fbe40a — Scripts to generate or retrieve training

data

What: Created shell scripts to automate dataset generation.

Why: Manual data collection doesn't scale. Needed 600+ samples for balanced training.

Commit 2: c843e16 — Wrote the function to calculate Shannon entropy

What: First implementation of `shanon_helper()`.

```
# BEFORE (Naive approach)
probabilities = {byte: count / total_bytes for byte, count in counter.items()}
for prob in probabilities.values():
    entropy += prob * math.log2(prob)
entropy = 0.0 - entropy # Negation at the end
```

Problem: Created an intermediate dictionary (`probabilities`) unnecessarily—wasted memory on large files.

Commit 3: 7b2fe64 — Optimized the `shanon_helper` function

What: Eliminated the intermediate dictionary and added safety checks.

```
# AFTER (Optimized)
for count in counter.values():
    prob = count / total_bytes
    entropy -= prob * math.log2(prob) # Direct subtraction
```

Improvement	Before	After
Memory allocation	Dict of 256 entries	None (direct iteration)
Empty file handling	✗ Division by zero crash	✓ Returns 0.0 safely
Readability	Two-step (calculate, then negate)	One-step (subtract directly)

Commit 4: 1b71e0c — Updated extractor.py and generate_custom_malware.sh

What: Two critical changes.

Change A: Refactored shanon_helper to accept binary data (not file path)

```
# BEFORE: Tightly coupled to file I/O
def shanon_helper(filepath):
    with open(filepath, 'rb') as file:
        file_bytes = file.read()

# AFTER: Pure function, accepts any bytes
def shanon_helper(binary_data):
    counter.update(binary_data)
```

Improvement	Impact
Testability	Can now unit test with <code>b'\x00\xff'</code> without creating files
Reusability	Same function works for PE sections, raw buffers, network streams
Performance	Avoids redundant file opens when PE is already parsed

Change B: Fixed the "Hex Trap" in malware generation

The Bug: Generating payloads as hex strings (A-F0-9) limited entropy to ~4.0 bits.

Hex uses only 16 characters (4 bits of entropy max), not 256 possible byte values (8 bits).

```
# BEFORE (Limited entropy)
PAYLOAD_DATA=$(head /dev/urandom | tr -dc A-F0-9 | head -c $PAYLOAD_SIZE)
const char encrypted_payload[] = "$PAYLOAD_DATA"; # Stored as ASCII string

# AFTER (True randomness)
PAYLOAD_ARRAY=$(head -c $PAYLOAD_SIZE /dev/urandom | od -An -v -t x1 | ...)
unsigned char encrypted_payload[] = { $PAYLOAD_ARRAY }; # Stored as raw bytes
```

Metric	Before (Hex String)	After (Raw Bytes)
Max theoretical entropy	4.0 bits	8.0 bits

Metric	Before (Hex String)	After (Raw Bytes)
Actual measured entropy	~3.8 bits	~7.99 bits
Realism	No real malware uses hex	Mimics AES/RC4 encrypted payloads

Commit 5: a1dd06c — Orchestrator v1 and get_structural_features

What: Introduced the Orchestrator pattern and added structural features.

New Features Extracted

Feature	Why It Matters
num_sections	Packed malware often has unusual section counts
virtual_size_ratio	Ratio > 1.0 indicates runtime unpacking (malware signal)
raw_size vs virtual_size	Discrepancy reveals hidden payloads

Architectural Change: Orchestrator Pattern

```
def extract_all_features(filepath) -> dict: # Orchestrator
    pe = pefile.PE(filepath)
    features = []
    features.update(get_entropy_features(pe))      # Worker 1
    features.update(get_structural_features(pe))   # Worker 2
    pe.close()
    return features
```

Benefit	Explanation
Extensibility	Add <code>get_import_features()</code> without touching entropy logic
Single Responsibility	Each worker does one thing well
Resource Safety	<code>pe.close()</code> in one place (now in <code>finally</code> block)

Commit 6: builder.py — Dataset CSV Generator (Bug Fixes)

What: Created `builder.py` to batch-process all PE files and generate a training CSV.

⚠ Bug A: `os.walk` Yields a Tuple, Not a File Path

The Mistake:

```
for file in os.walk('../data/benign', topdown=True):
    row = extract_all_features(file) # ❌ CRASH: passing a tuple!
```

What Happened:

`os.walk()` doesn't yield file paths—it yields a **3-tuple**:

`(current_folder, list_of_subfolders, list_of_files)`. The code passed the entire tuple to `pefile.PE()`, which expects a string path.

The Fix:

```
for root, dirs, files in os.walk('../data/benign'): # ✓ Unpack the tuple
    for filename in files:
        full_path = os.path.join(root, filename) # ✓ Reconstruct full path
        row = extract_all_features(full_path)
```

Component	What It Contains	Example
root	Current directory path	../data/benign/subdir
dirs	List of subdirectory names	['folder1', 'folder2']
files	List of file names (no path!)	['file1.exe', 'file2.dll']

⚠ Bug B: Memory Trap — Accumulating All Rows Before Writing

The Mistake:

```

rows = []
for root, dirs, files in os.walk('../data'):
    for filename in files:
        row = extract_all_features(...)
        rows.append(row) # ❌ Stores everything in RAM

# Write all at once at the end
writer.writerows(rows)

```

The Problem:

At small scale (600 files), this works. At production scale (1 million files), storing 1 million dictionaries in RAM before writing a single line causes **memory exhaustion**.

The Fix — Streaming Pattern:

```

with open('dataset.csv', 'w') as csvfile:
    writer = csv.DictWriter(csvfile, fieldnames=fieldnames)
    writer.writeheader()

    for root, dirs, files in os.walk('../data'):
        for filename in files:
            row = extract_all_features(...)
            if row:
                writer.writerow(row) # ✅ Write immediately, then discard

```

Approach	Memory Usage (1M files)	Disk I/O
Accumulate then write	~8GB+ (crash likely)	One burst at end
Stream (write immediately)	~constant (~1KB)	Continuous, efficient

Lesson Learned:

Never accumulate unbounded data in memory. Write/process incrementally ("streaming") to keep RAM usage constant regardless of input size.

⚠ Bug C: Missing File Extension Filter

The Mistake:

```

for filename in files:
    row = extract_all_features(os.path.join(root, filename))
    # ❌ Will try to parse .sh, .DS_Store, .txt, etc.

```

The Problem:

The data directory might contain non-PE files (shell scripts, macOS metadata, logs). Attempting to parse them with `pefile` causes unnecessary errors and clutters logs.

The Fix:

```

for filename in files:
    if filename.endswith('.exe', '.dll')): # ✅ Filter before processing
        row = extract_all_features(os.path.join(root, filename))

```

File Type	Before (No Filter)	After (With Filter)
malware.exe	✓ Processed	✓ Processed
generate.sh	✗ Error logged	▶ Skipped silently
.DS_Store	✗ Error logged	▶ Skipped silently

Best Practice:

Always validate input before processing. Filtering early ("fail fast") prevents wasted computation and keeps logs clean.

4. Key Technical Decisions

Decision 1: VirtualSize vs. RawSize for Entropy Calculation

Observation: Malware often has `VirtualSize > SizeOfRawData` because it unpacks at runtime.

Problem: `section.get_data()` returns raw bytes padded with nulls, diluting entropy.

Solution:

```
section_data_trimmed = section.get_data()[:section.Misc_VirtualSize]
```

Scenario	Without Trimming	With Trimming

Scenario	Without Trimming	With Trimming
10KB encrypted payload + 90KB null padding	~0.8 entropy	~7.9 entropy

Decision 2: Statistical Aggregation for Variable-Length Data

Problem: ML models need fixed-length input vectors, but PE files have 3-10+ sections.

Solution: Summarize with statistical moments:

```
entropy_summary = {
    'avg_entropy': mean(entropies),      # Overall complexity
    'max_entropy': max(entropies),       # THE KEY SIGNAL (encrypted section)
    'min_entropy': min(entropies),       # Null-padded sections
    'std_entropy': np.std(entropies)     # Variance between sections
}
```

Why std_entropy is crucial:

Benign files have *consistent* entropy across sections. Malware has *spikes* (encrypted payload vs. empty padding). Standard deviation captures this variance.

Decision 3: Error Handling at Orchestrator Level

Pattern: Workers assume valid input; Orchestrator handles exceptions.

⚠ Bug Discovered: UnboundLocalError in finally Block

The Mistake:

```
def extract_all_features(filepath) -> dict:
    try:
        pe = pefile.PE(filepath) # If this fails, pe is never assigned
        # ... call workers ...
    except pefile.PEFormatError as e:
        print(f"Error processing {filepath}: {e}")
        return None
    finally:
        pe.close() # ❌ CRASH: UnboundLocalError if file doesn't exist
```

What Happened:

When a `FileNotFoundException` occurred (wrong file path), the `pe = pefile.PE(filepath)` line threw an exception *before* `pe` was assigned. The `finally` block still executed, but `pe` didn't exist—causing a *second* error that masked the real problem.

The Fix:

```
def extract_all_features(filepath) -> dict:
    pe = None # ✅ Initialize before try block
    try:
        pe = pefile.PE(filepath)
        # ... call workers ...
    except pefile.PEFormatError as e:
        print(f"Error processing {filepath}: {e}")
        return None
    finally:
        if pe is not None: # ✅ Guard before closing
            pe.close()
```

Issue	Before	After
Missing file	UnboundLocalError masks real error	Shows actual FileNotFoundError
Invalid PE	UnboundLocalError	Graceful None return
Valid PE	Works	Works

Lesson Learned:

Always initialize resources to `None` before a `try` block when using `finally` for cleanup. This pattern is called "**Initialize-Try-Finally**" and prevents cascading errors.

Benefit	Impact
Cleaner workers	No redundant try/except in every function
Consistent error format	All errors logged the same way
Memory safety	<code>finally</code> ensures handles are closed even on crash
Error transparency	Real exceptions aren't masked by cleanup failures

5. Pythonic Implementation Details

5.1 Counter.update() vs. Manual Loop

```
# ❌ WRONG: Treats entire bytes object as one key
counter[file_bytes] += 1 # Result: Counter({b'...': 1})

# ❌ SLOW: Correct but inefficient
for byte in file_bytes:
    counter[byte] += 1

# ✅ OPTIMAL: C-optimized iteration
counter.update(file_bytes) # ~10x faster for large files
```

Why update() is faster: Implemented in CPython's C internals, avoiding Python's interpreter overhead per byte.

5.2 Binary Data Representation in Python

Misconception: Expecting `b'\x7f\x45\x4c\x46'` when printing bytes.

Reality: Python displays printable ASCII: `b'\x7fELF'`

Insight: Internally, Python treats bytes as integers 0-255. To see them:

```
>>> list(b'ELF')
[69, 76, 70]
```

This is exactly what Shannon entropy math uses—256 possible byte values.

6. Architecture & Code Organization

```
extractor.py
├── shanon_helper(binary_data)      # Pure function: bytes → entropy float
├── get_entropy_features(pe)        # Worker: PE → entropy dict
├── get_structural_features(pe)     # Worker: PE → structure dict
└── extract_all_features(filepath)   # Orchestrator: path → final feature dict
```

Design Principles Applied

Principle	Implementation
Single Responsibility	Each function does exactly one thing
Dependency Injection	Workers receive <code>pe</code> object, don't open files
Fail Fast	Orchestrator validates PE before calling workers
Resource Management	<code>try/finally</code> ensures <code>pe.close()</code> always runs

7. Model Training & Results

7.1 Training Pipeline (`trainer.py v1`)

```
# Load and prepare data
df = pd.read_csv('.../data/dataset.csv')

# Drop multicollinear and non-predictive features
X = df.drop(columns=['is_malicious', 'raw_size', 'virtual_size', 'filename'])
y = df['is_malicious']

# Stratified split (maintains class balance in train/test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Train Random Forest
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```


Why These Features Were Dropped

Dropped Feature	Reason
<code>filename</code>	Non-predictive (string identifier, not a signal)
<code>raw_size</code>	Multicollinear with <code>total_file_size</code>
<code>virtual_size</code>	Multicollinear with <code>raw_size</code> ; ratio captures the relationship

7.2 Exploratory Data Analysis (EDA)


Before training, we visualized the dataset to understand feature relationships and validate our assumptions.


Feature Correlation with Target (`is_malicious`)



Insight: Shows which features have the strongest correlation with malware classification. Entropy features (`max_entropy` , `std_entropy`) show the highest positive correlation with `is_malicious` .


Max Entropy Distribution by Class





Insight: Clear separation between benign (class 0) and malicious (class 1) samples. Malware consistently shows higher `max_entropy` values (~7.5-8.0), while benign files cluster around 5.0-6.5. This visual confirms `max_entropy` is the strongest discriminator.

Multicollinearity Between Features



Insight: Reveals which features are redundant:

- `raw_size` and `virtual_size` are highly correlated (~0.95+) → dropped `raw_size`
- `virtual_size` and `total_file_size` overlap → kept `total_file_size`, dropped `virtual_size`
- `virtual_size_ratio` captures the relationship without redundancy

Why this matters: Multicollinear features don't add new information but can destabilize model

coefficients and inflate feature importance scores.

7.3 Results: Perfect Classification

Confusion Matrix


		Predicted
		0 1
Actual	0	[99 0]
	1	[0 100]

Metric	Value
True Negatives (Benign correct)	99
False Positives	0
False Negatives	0
True Positives (Malware correct)	100

Classification Report

Class	Precision	Recall	F1-Score	Support
0 (Benign)	1.00	1.00	1.00	99
1 (Malware)	1.00	1.00	1.00	100
Accuracy			1.00	199

7.4 Feature Importance Analysis (The "Why")



Interpretation

Rank	Feature	Importance	What It Tells Us
1	max_entropy	32.75%	The smoking gun — encrypted payloads hit ~7.99 entropy
2	std_entropy	25.16%	Malware has <i>variance</i> (encrypted section vs. null padding)
3	avg_entropy	19.43%	Overall file complexity
4	min_entropy	11.49%	Benign files rarely have near-zero sections
5	size_of_headers	6.12%	Anomalous headers can indicate tampering
6-8	Others	< 3% each	Minor contributors

Key Insight: Entropy Features Dominate

The top 4 features are **all entropy-based**, accounting for **88.83%** of the model's decision-making power. This validates the project's core hypothesis:

Section-aware entropy analysis is the strongest signal for detecting packed/encrypted malware.

7.5 Caveats & Next Steps

⚠ Why 100% accuracy might be misleading:

Concern	Explanation
Synthetic data	Custom malware was generated with intentionally high entropy—real malware is more varied

Concern	Explanation
Small dataset	199 test samples; production needs 10,000+
No adversarial samples	Sophisticated malware uses entropy masking techniques

Recommended next steps:

1. Test on real-world malware samples (VirusTotal, MalwareBazaar)
2. Add adversarial samples (entropy-normalized malware)
3. Cross-validate with k-fold to ensure robustness

8. Future Improvements

- Import Address Table (IAT) Analysis** — Detect suspicious API calls (`VirtualAlloc` , `WriteProcessMemory`)
- Section Name Analysis** — Flag unusual names like `UPX0` , `.enigma`
- String Extraction** — Find hardcoded C2 URLs or registry keys
- YARA Rule Integration** — Cross-reference with known malware signatures
- Batch Processing Pipeline** — Parallel extraction with `multiprocessing`
- Train baseline model** — Random Forest achieving 100% on synthetic dataset
- Real-world validation** — Test on MalwareBazaar/VirusTotal samples

Appendix: Quick Reference

Shannon Entropy Formula

$$H(X) = - \sum_{i=0}^{255} p(x_i) \log_2 p(x_i)$$

Where $p(x_i)$ is the probability of byte value i appearing in the data.

Entropy Interpretation

Entropy Range	Meaning	Example
---------------	---------	---------

Entropy Range	Meaning	Example
0.0 - 1.0	Highly structured/repetitive	Null-padded sections
4.0 - 5.0	Plain text / code	.text section
7.0 - 8.0	Encrypted / compressed	AES payload, packed UPX

Document maintained as part of the EntropyX malware classification project.